

# Microservices


The three principals of the microservices architecture style.

- \* Microservices are ideal for big systems
- \* Microservice architecture is goal-oriented not solution-oriented
- \* Microservices are focused on replaceability

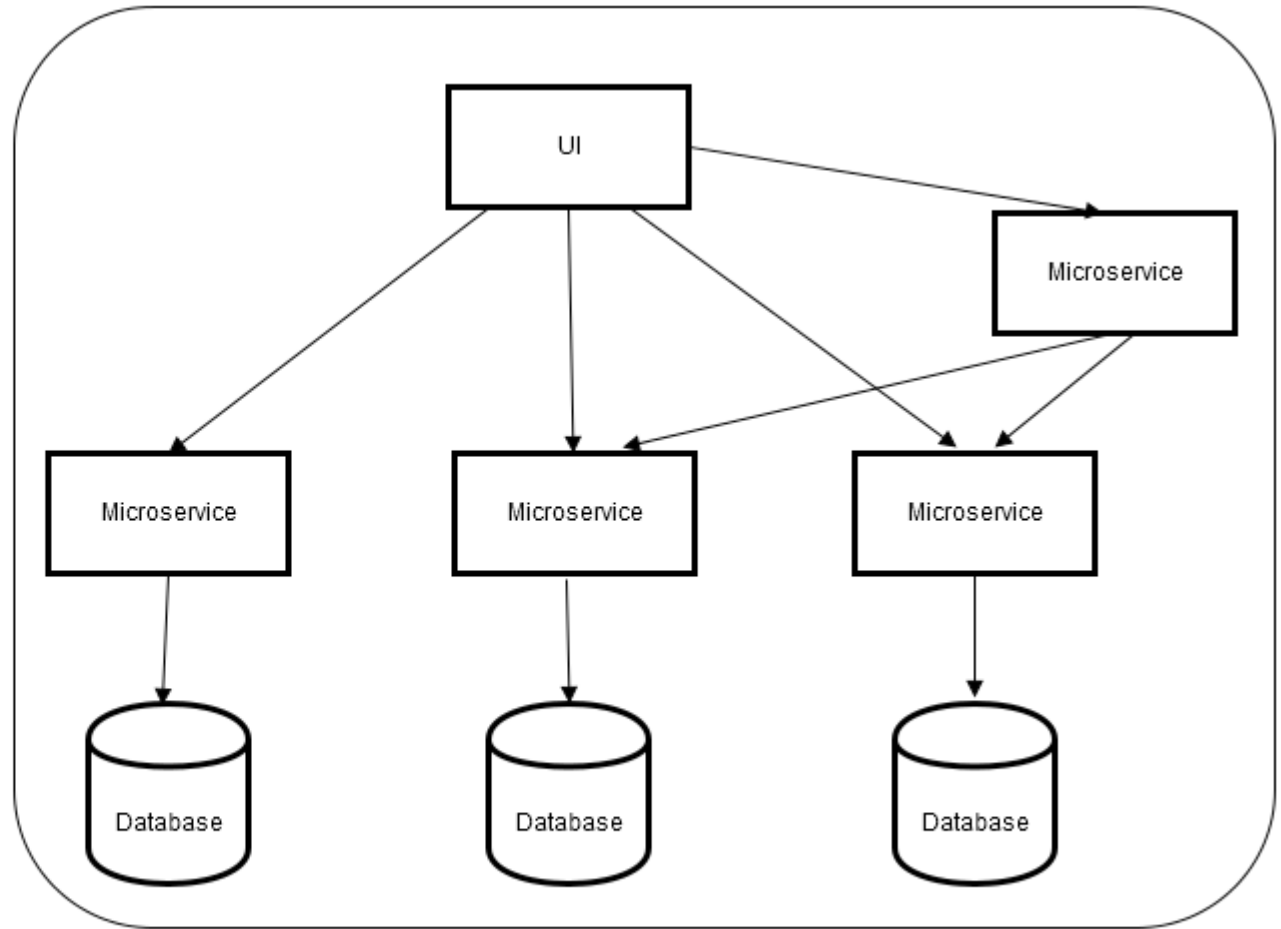
# microservice applications share some important characteristics

- Small in size
- Messaging enabled
- Bounded by contexts
- Autonomously developed
- Independently deployable
- Decentralized
- Built and released with automated processes

# Web Service vs Microservice



Monolithic Architecture



Microservices Architecture

# Web Service

Web Service is a way to expose the functionality of an application to other application, without a user interface. It is a service which exposes an API over HTTP.

Web Services allow applications developed in different technologies to communicate with each other through a common format like XML, Json, etc.

# Microservices

Micro Service is independently deployable service modeled around a business domain. It is a method of breaking large software applications into loosely coupled modules, in which each service runs a unique process and communicates through APIs. It can be developed using messaging or event-driven APIs, or using non-HTTP backed RPC mechanisms.

# Difference

Microservice: Non-Java EE based that can address cross-cutting concerns in request/response chain.

Webservice: Java EE based that cannot manipulate immutable request and response. All cross-cutting concerns must be addressed at network level. Commercial gateway.

# Difference

- Microservice: low latency, high throughput and small memory footprint.
- Webservice: high latency, low throughput and big memory footprint.

### Best plaintext responses per second, Dell servers at ServerCentral (233 tests)

Framework	Best performance (higher is better)	Cls	Lng	Plt	FE	Aos	IA	Errors
octane	4,366,106   <div></div> 100.0%	Plt	C	Non	Non	Lin	Rea	126
rapidoid-http-fast	3,739,042   <div></div> 85.6%	Plt	Jav	Rap	Non	Lin	Rea	0
ulib	3,731,388   <div></div> 85.5%	Plt	C++	Non	ULi	Lin	Rea	1
rapidoid	3,689,000   <div></div> 84.5%	Plt	Jav	Rap	Non	Lin	Rea	0
tokio-minihttp	3,538,853   <div></div> 81.1%	Mcr	Rus	Rus	tok	Lin	Rea	0
libreactor	3,463,733   <div></div> 79.3%	Plt	C	lib	Non	Lin	Rea	0
colossus	3,055,275   <div></div> 70.0%	Mcr	Sca	Akk	Non	Lin	Rea	0
light-java	2,834,876   <div></div> 64.9%	Plt	Jav	Lig	Non	Lin	Rea	0

### Best database-access responses per second, single query, Dell servers at ServerCentral (267 tests)

Framework	Best performance (higher is better)	Cls	Lng	Plt	FE	Aos	DB	Dos	Orm	IA	Errors
ulib-mongodb	210,393   <div></div> 100.0%	Plt	C++	Non	ULi	Lin	Mo	Lin	Mcr	Rea	0
h2o	185,682   <div></div> 88.3%	Plt	C	Non	Non	Lin	Pg	Lin	Raw	Rea	0
ulib-postgres	179,031   <div></div> 85.1%	Plt	C++	Non	ULi	Lin	Pg	Lin	Mcr	Rea	0
cpoll_cppsp-postgres	177,089   <div></div> 84.2%	Plt	C++	Non	Non	Lin	Pg	Lin	Raw	Rea	0
light-java	176,098   <div></div> 83.7%	Plt	Jav	Lig	Non	Lin	Pg	Lin	Raw	Rea	0

### Best fortunes responses per second, Dell servers at ServerCentral (238 tests)

Framework	Best performance (higher is better)	Cls	Lng	Plt	FE	Aos	DB	Dos	Orm	IA	Errors
ulib-postgres	180,731   <div></div> 100.0%	Plt	C++	Non	ULi	Lin	Pg	Lin	Mcr	Rea	0
ulib-mongodb	180,574   <div></div> 99.9%	Plt	C++	Non	ULi	Lin	Mo	Lin	Mcr	Rea	0
gemini-postgres	177,682   <div></div> 98.3%	Ful	Jav	Svt	Res	Lin	Pg	Lin	Mcr	Rea	0
cutelyst-thread-pg-e	164,422   <div></div> 91.0%	Plt	C++	Qt	Non	Lin	Pg	Lin	Raw	Rea	0
cutelyst-thread-pg-r	164,302   <div></div> 90.9%	Plt	C++	Qt	Non	Lin	Pg	Lin	Raw	Rea	0
ur/web	155,817   <div></div> 86.2%	Ful	Ur	Ur/	Non	Lin	Pg	Lin	Mcr	Rea	0
light-java	154,086   <div></div> 85.3%	Plt	Jav	Lig	Non	Lin	Pg	Lin	Raw	Rea	0



# Microservice Security Is Harder

- Win

Every service only has access to what it needs to perform its function

- Lose

Much larger attack surface(internal threats)

How do other services know who's accessing them?

How can other services trust each other?

# OAuth2

## Delegated Authorization

- A protocol for conveying authorization decisions via a token
- Standard means of obtaining a token (aka 4 OAuth2 grant types)

Authorization Code

Resource Owner Password Grant

Implicit

Client Credentials

- Users and Clients are separate entities

I am authorizing this app to perform these actions on my behalf

# What is OAuth2 Not

## OAuth2 is not Authentication

- The user must be authenticated to obtain a token
- How the user is authenticated is outside of the spec
- How the token is validated is outside the spec
- What the token contains is outside the spec.

# What is OpenID Connect

## Delegated Authentication

- A protocol for conveying user identity via a signed JWT
- Built on top of OAuth2
- Standard means of obtaining an ID token
- Standard means of verifying ID token
- Steve is authorizing this app to perform these actions on his behalf and here is his email and role in case you need it

# Beyond OAuth2: End to End Microservice Security

- Token Propagation
- ID token and Client Credentials token
- New Tokens via Token Exchange
- Data Integrity
- Data Confidentiality

# Token Propagation

- The token is too powerful
- Can be used to do anything to the system as that user until it expires
- Token leakage is a big deal
- Internal fraud is easy

# ID token and CC token

- A program fooled into misusing its authority  
e.g. when one app fully trusts another by virtue of the app's identity
- A and B fully trust Resource

Where's the proof that Resource is acting faithfully?

How can A and B know if the User is actually authorized?

- If Resource is compromised, A and B will be compromised

# Total trust in the bank

- If apps fully trust one another, do teams as well?

Transitively, perhaps

Do banks really work that way?

- What happens when you have no trust boundaries?

You don't check the other person's work

You allow a single person to perform multiple critical tasks

You have no separation of duties



# Insiders

- Most attackers are insiders. Over 60%.

# Confused deputy mitigations

- Authorize based on more than just caller's identity

The user and their scopes

- Send both client and user's token

Still vulnerable, the combination is not integrity protected

- Authorize based on a composite token
- Authorize based on a call stack

No information loss

Specify allowable behaviours with high precision

# New Tokens via Token Exchange

- Given Actor + Subject + Audience, get a new token

Policy decision given caller, user and intent

New token expresses caller, user and intent

- Given Actor + previous token + Audience, get a new token

Policy decision based on delegation chain(call stack)

Policy decision based on aggregated trust

# Token Exchange

- Pros

User, client and call stack

Narrow audience and scope

Trust boundaries are unambiguous

Centralized policy management and de-centralized policy enforcement

- Cons

Network and AS overhead

Security vs Performance

Policy Management vs Agility

# HTTPS doesn't solve message security

- Does Cart really need to see the CC info?

Payment is the only service that really use it

- Only Point to Point confidentiality and integrity are assured
- Rearrange services so Payment can be directly called by Shop

Limits architectural choice

Overcomplicates Shop

# End to End Message Security

- Cannot be solved at the network layer
- Necessarily an application layer concern
- Stop assuming trust based on solely on caller's identity
- Limit the effect of token leakage and misuse
- Break out of the performance vs security tradeoff

# JOSE for Message Security

- Use public/private key pairs
- Sender signs message with private key

Integrity and non-repudiation

- Sender encrypts signed messages with public key of recipient

Confidentiality

- Recipient decrypts with its private key, verifies with sender's public key

# Key Pairs

- Public Key Infrastructure(PKI)
- Certificate Authority
- OAuth2 Server



# Oauth2 and Service Discovery

- Oauth2 server generates key pairs
- Service first time start to download key pair with `client_id` and `client_secret`
- Service register host, port and public key to service discovery (consul)
- Other services can find public key of other service from discovery

# Sign and Verify Data

- The sender signs with its private key

HTTP request, response, headers

Messaging body

- Receiver users the registry to get sender's public key
- Use the public key to verify payloads
- Payloads are traceable to the individual service that registered

# Encrypt and Decrypt Data

- Sender uses the registry to get receiver's public key

- Use public key to encrypt payloads

sign, then encrypt

multiple receivers possible via JWE JSON  
Serialization

encrypt entire bodies, objects or single field

- Use it when needed

# Distributed OAuth2

- End to End message confidentiality with JWE
- Message integrity via JWS

Service authentication for free

- Still need to deal with authorization

Assert and authorize call stack

Limit the effect of token leakage and misuse

Performance vs security tradeoff

# Self Issued JWT

- Services authenticate with JWTs they create and sign themselves

JWT is just a specific use case of JWS

Services authenticate by verifying with sender's public key

- Create as many as you need without network overhead

# Signle Use JWT

- Short expiry
- JTI can prevent replay attacks
- Unambiguous intent

Express the intended recipient and operation to be performed

- Very limited power

Can only be used for intended purpose and only once

- Services only accept these JWTs

Reject the bear user token as it has no power on its own

# Nested, Self Issued JWT

- Incoming JWTs can be nested in another JWT for use downstream

Call stack expressed as nested JWTs

- Verifiable chain of custody

Like an audit trail

Call stack verified by verifying each JWT recursively

- Unbreakable chain of custody

Sender and receiver is encoded in the chain

Chains are wrapped in another JWS, with doesn't get propagated

Chain truncation can be detected

# Externalized Policy, Embedded Decision Making

- Services make authorization decisions on their own
- Policy as externalized configuration

Not hard-coded into services

- No additional calls required for authorization

Good performance

- Policy can be flexible

Allow services to evolve